

L6136 - DEEP LEARNING ON GPUS: FROM LARGE SCALE TRAINING TO EMBEDDED DEPLOYMENT

Julie Bernauer, 4/5/2016

AGENDA

0. Introduction
1. Connecting to the Amazon cloud
2. Train an image classification model using Digits
3. Caffe on the TX1: the fp16 whitepaper
4. Deploy the image classification model on TX1
5. Fun with the webcam

INTRODUCTION

ACCELERATING INSIGHTS

“Now You Can Build Google’s
\$1M Artificial Brain on the Cheap”



GOOGLE DATACENTER



1,000 CPU Servers
2,000 CPUs • 16,000 cores

600 kWatts
\$5,000,000


STANFORD AI LAB



3 GPU-Accelerated Servers
12 GPUs • 18,432 cores

4 kWatts
\$33,000

MODERN AI



2012: GOOGLE BRAIN





2016: AlphaGO

WHAT IS DEEP LEARNING?



Input



Result







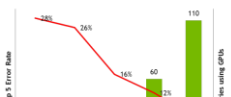

IMAGE CLASSIFICATION

Deep Learning Framework



NVIDIA DEEP LEARNING SDK

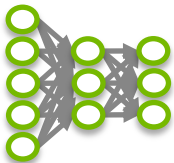

High Performance GPU-Acceleration for Deep Learning




Image Classification Object Detection
COMPUTER VISION
Voice Recognition

Language Translation

SPEECH AND AUDIO
Recommendation Engines


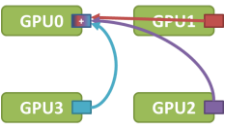
Sentiment Analysis

NATURAL LANGUAGE PROCESSING
DEEP LEARNING FRAMEWORKS
cuDNN
DEEP LEARNING
cuBLAS

cuSPARSE

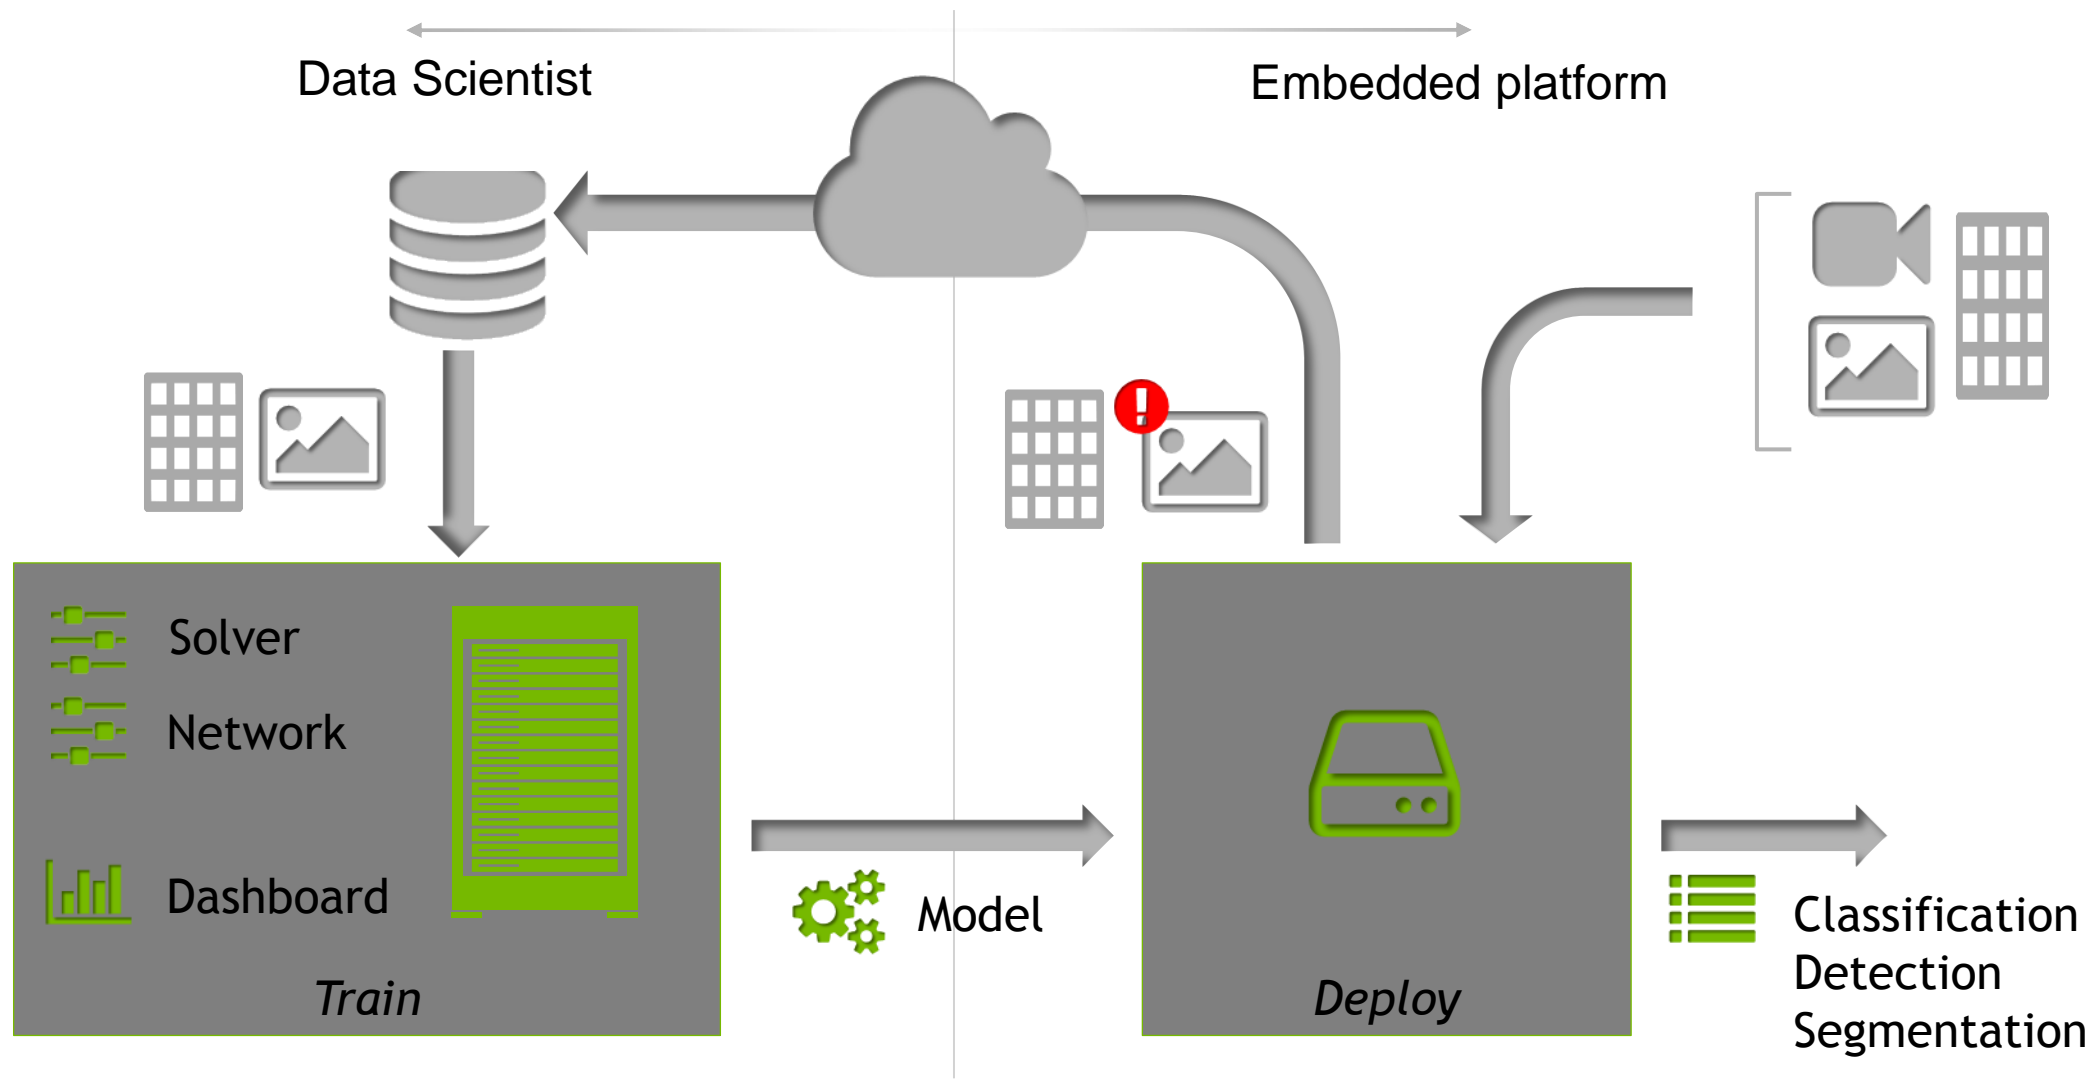


cuFFT

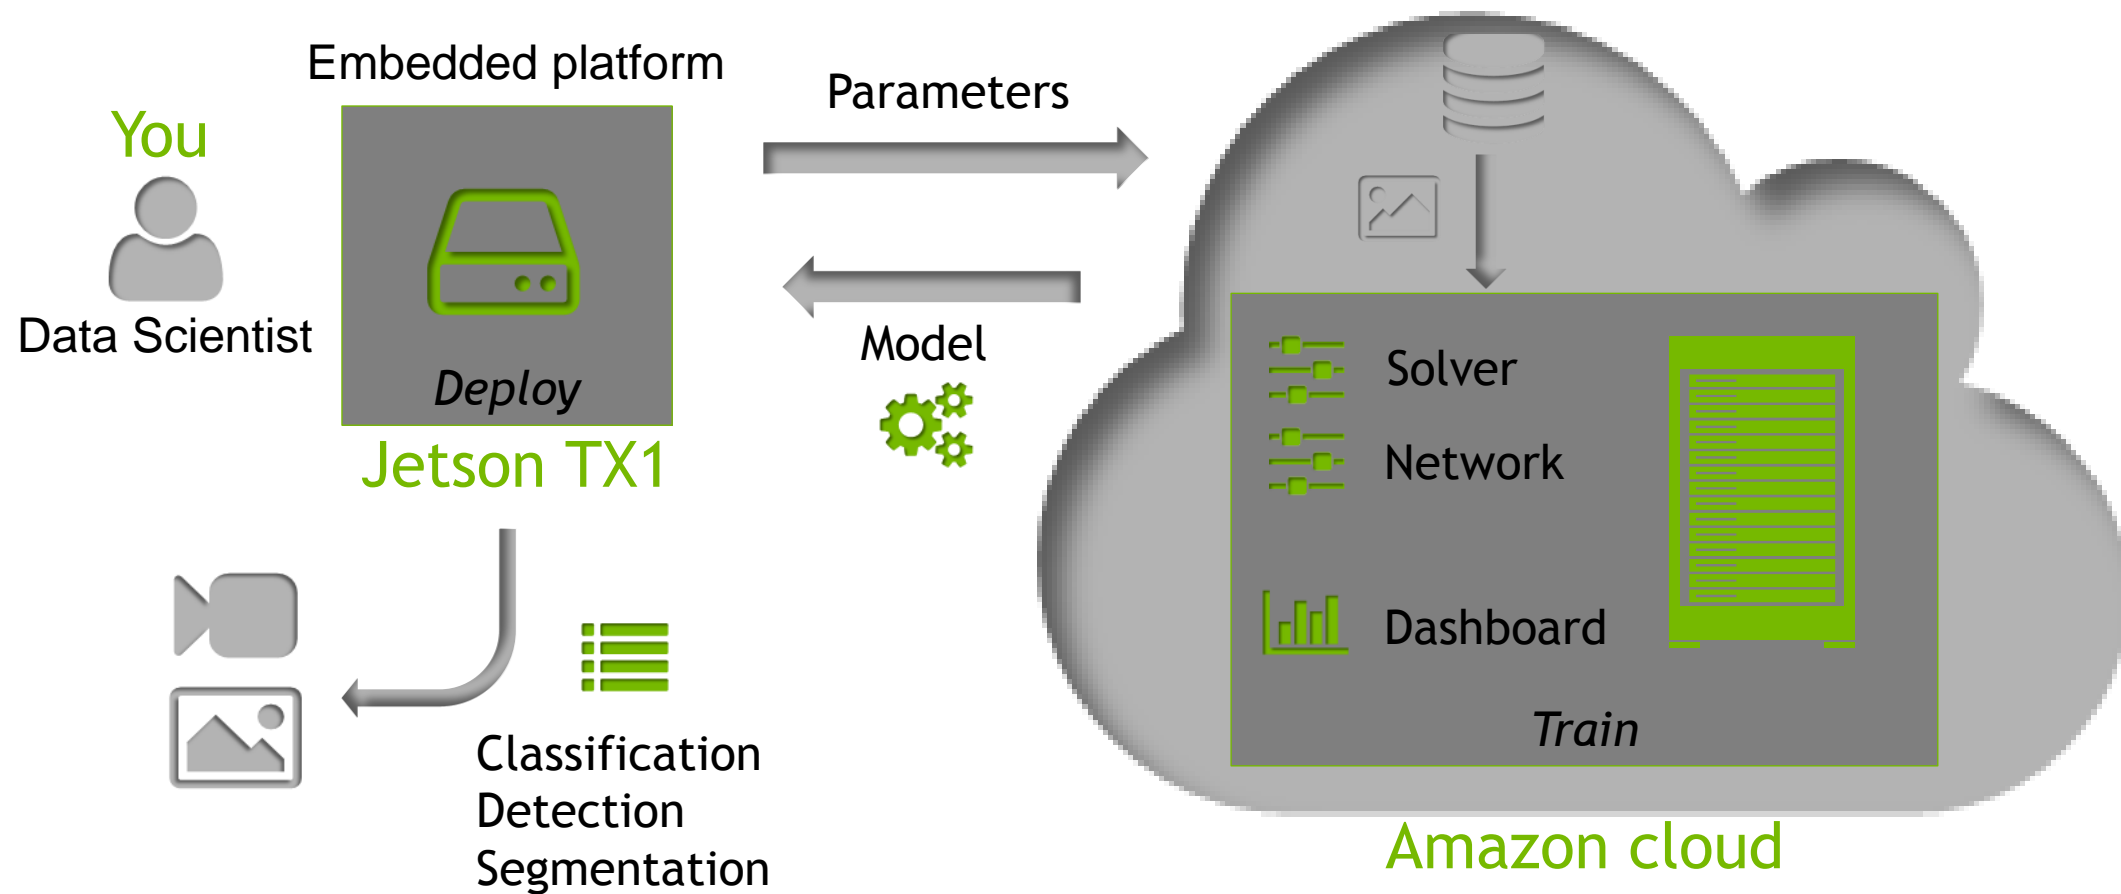
**MATH LIBRARIES**
NCCL
MULTI-GPU

TODAY

DEEP LEARNING ON GPUS: AN END-TO-END SOLUTION




TODAY'S LAB



EMBEDDED DEPLOYMENT

On Jetson TX1



Inference at 258 img/s

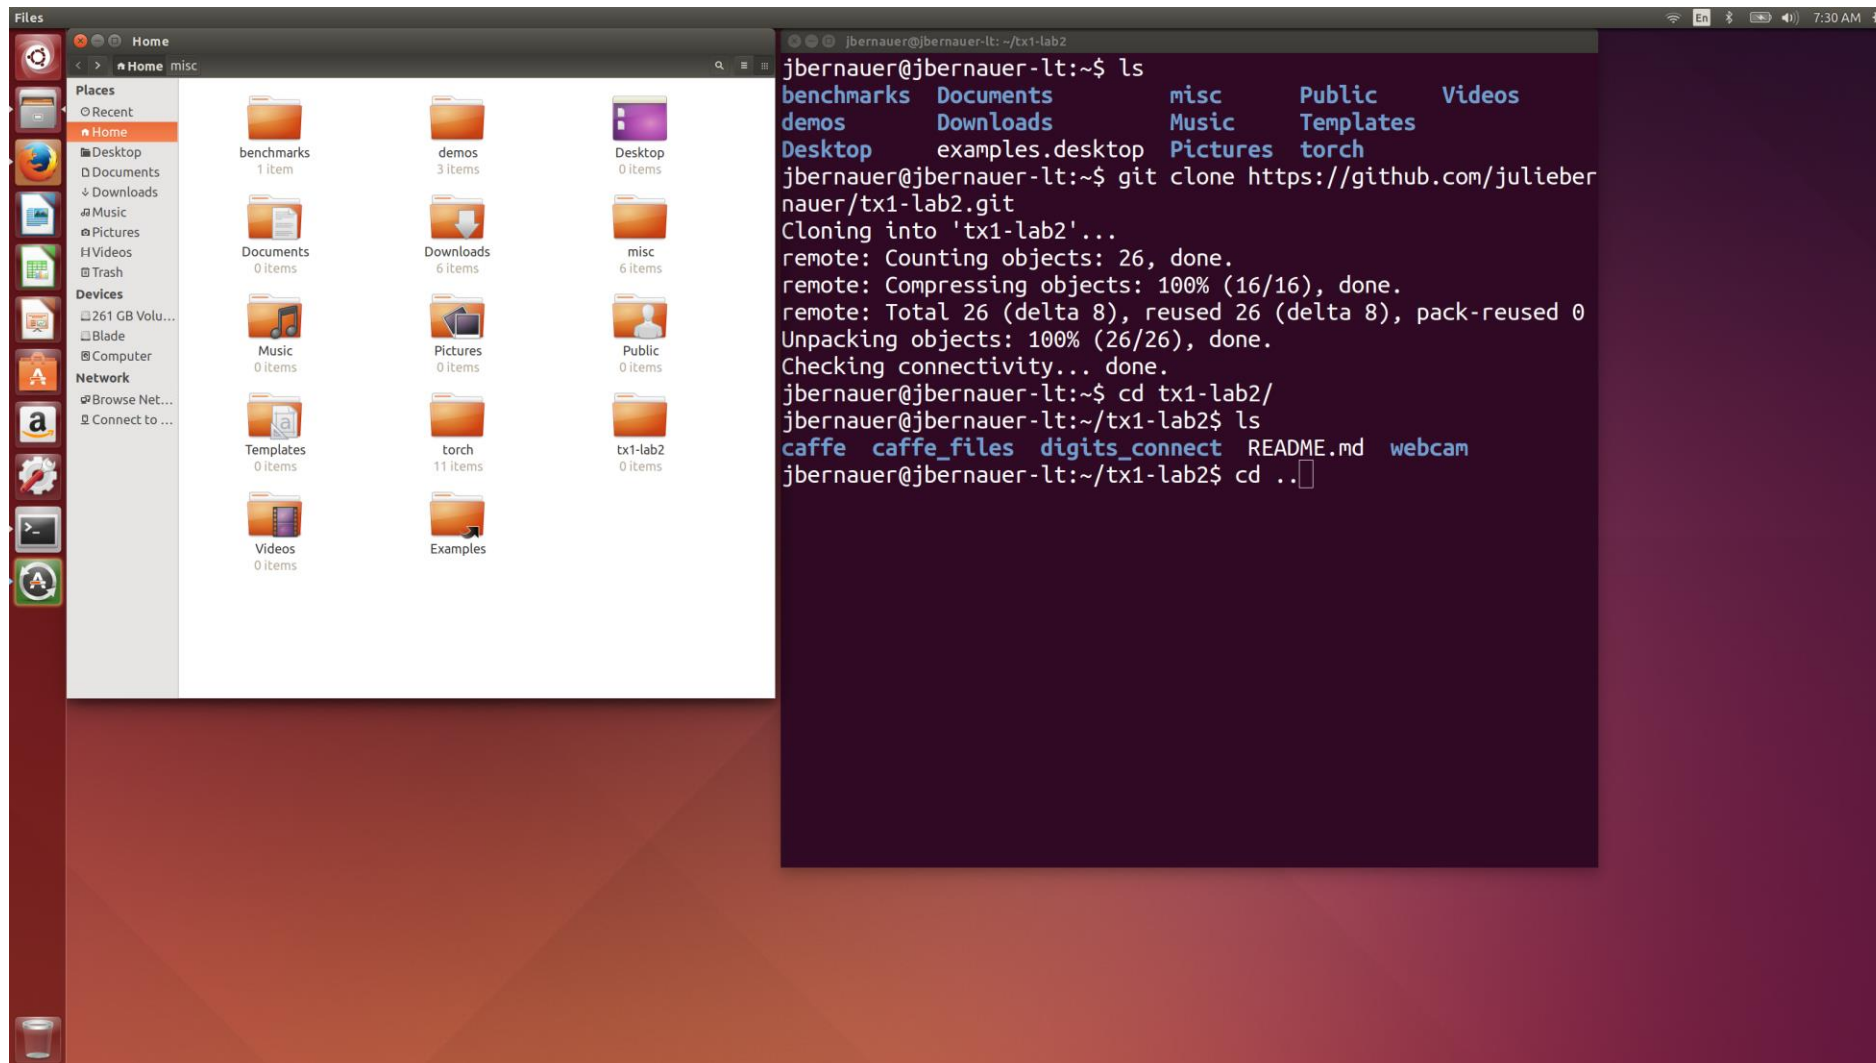
No need to change code

Simply compile Caffe and copy a
trained `.caffemodel` to TX1

SET UP

USING THE TERMINAL

Linux Ubuntu 14.04



0- DOWNLOADING LAB DATA

From the TX1

- Install git on the tx1: *sudo apt-get install git*
- Data for the TX1 is on github, in the terminal, type:
git clone <https://github.com/juliebernauer/tx1-lab2.git>
- The instructions can also be seen on the github page:
<https://github.com/juliebernauer/tx1-lab2>

1- CONNECTING TO THE AMAZON CLOUD

From the TX1

- Connect to <https://nvlabs.qwiklabs.com>
- Pick the *Digits and Jetson TX1* lab, click the *Start* button.
- Wait for the lab to launch
- Click on *Connection -> Start your lab*
- Click on the DIGITS link

GETTING STARTED WITH DEEP LEARNING DIGITS

2- TRAIN AN IMAGE CLASSIFICATION MODEL

Using digits on Amazon EC2

In the web browser, connect to your amazon instance DIGITS web server

DIGITS

LoginInfo

Home

2/2 GPUs available

No Jobs Running

Datasets

New Dataset

Images

Models

View details

New Model

Images

CREATE YOUR USERNAME

DIGITS

[Login](#)

[Info](#) ▾

Login

Username ?

Submit

CREATE YOUR DATASET

DIGITS

Login

Info ▾

Home

2/2 GPUs available

No Jobs Running

Datasets

New Dataset

Images ▾

Classification

Other

Models [View details](#)

New Model

Images ▾

CREATE A DATASET

DIGITS

New Dataset

ubuntu (Logout)

Info

Image Type ?

Color

Image size ?

28 x 28

Resize Transformation ?

Squash

See example

Use Image Folder Use Text Files

Training Images ?

folder or URL

Minimum samples per class ?

2

Maximum samples per class ?

% for validation ?

25

% for testing ?

0

☐ Separate validation images folder

☐ Separate test images folder

/home/ubuntu/data/food

DB backend

LMDB

Image Encoding ?

PNG (lossless)

Dataset Name

Create a name

Create

EVALUATE YOUR DATASET

DIGITS

Image Classification Dataset

ubuntu (Logout)

Info ▾

Create DB (train)

Input File (before shuffling)

[train.txt](#)

DB Creation log file

[create_train_db.log](#)

DB Entries

9830

Category distribution

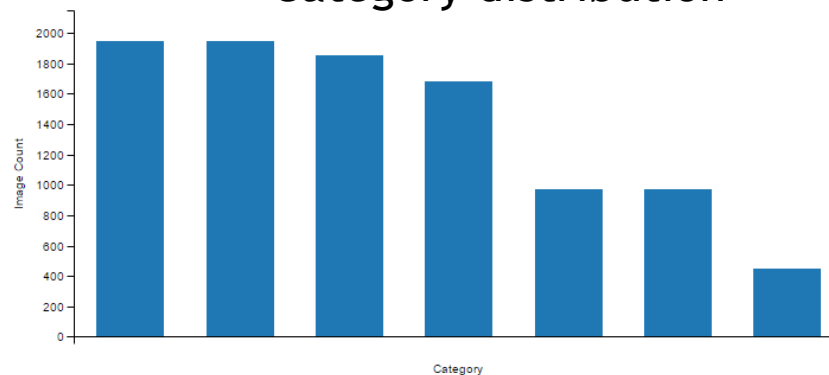


Image Mean:



Explore your
dataset

[Explore the db](#)

EXPLORE YOUR DATASET

Pick the *food_data* dataset


DIGITS
Image Classification Dataset
Login
Info

Exploring food_data (train_db) images


[Show all images](#) or filter by class:
[pizza](#)
[water_bottle](#)
[orange](#)
[lemon](#)
[cucumber](#)
[coffee_mug](#)
[background](#)
[bell_pepper](#)
[keyboard](#)
[pineapple](#)
[ipod](#)
[banana](#)
[strawberry](#)
[sunglasses](#)

Items per page: 10 - 25 - 50 - 100


«
0
1
2
3
4
5
6
...
2174
»




lemon




sunglasses




pineapple




bell_pepper




ipod




keyboard




keyboard




pineapple




banana



strawberry



cucumber



broccoli

TRAIN A NETWORK

DIGITS

New Model

ubuntu (Logout) Info

Select Dataset ⓘ

animals

☐ Use client side file

Python Layer File (server side) ⓘ

animals

Done 10:20:27 PM

Image Size
28x28

Image Type
COLOR

DB backend
lmdb

Create DB (train)
9830 images

Create DB (val)
3276 images

Data Transformations

Crop Size ⓘ
none

Subtract Mean ⓘ
Image

Solver Options

Training epochs ⓘ
30

Snapshot interval (in epochs) ⓘ
1

Validation interval (in epochs) ⓘ
1

Random seed ⓘ
[none]


Standard Networks

Previous Networks

Custom Network


Caffe

Torch (experimental)

Network	Details	Intended image size
 LeNet	Original paper [1998]	28x28 (gray)
<input type="radio"/> AlexNet	Original paper [2012]	256x256
<input type="radio"/> GoogLeNet	Original paper [2014]	256x256

REVIEW RESULTS

- What is your accuracy?
- How long did it take to train?



CLASSIFY AN IMAGE

DIGITS

Image Classification Model

Trained Models

Select Model

Epoch #30

Download Model

Test a single image

Image URL

Upload image

 No file chosen☒ Show visualizations and statistics ?

Test a list of images

Upload Image List

 No file chosen

Accepts a list of filenames or urls (you can use your val.txt file)

Image folder (optional)

Relative paths in the text file will be prepended with this value before reading

Number of images use from the file


Leave blank to use all

CLASSIFY AN IMAGE

Explore Network Response

DIGITS
animals
Classify One
Login
Info

Classify One Image
Clone Job
Delete Job

animals Image Classification Model


Predictions	
elephant	99.35%
domestic dog	0.31%
sea lion	0.19%
domestic cat	0.14%
tiger	0.01%

Job Status Done

- Initialized at 05:29:34 PM (1 second)
- Running at 05:29:35 PM (6 seconds)
- Done at 05:29:41 PM
- (Total - 7 seconds)

Infer Model Done

Notes
None

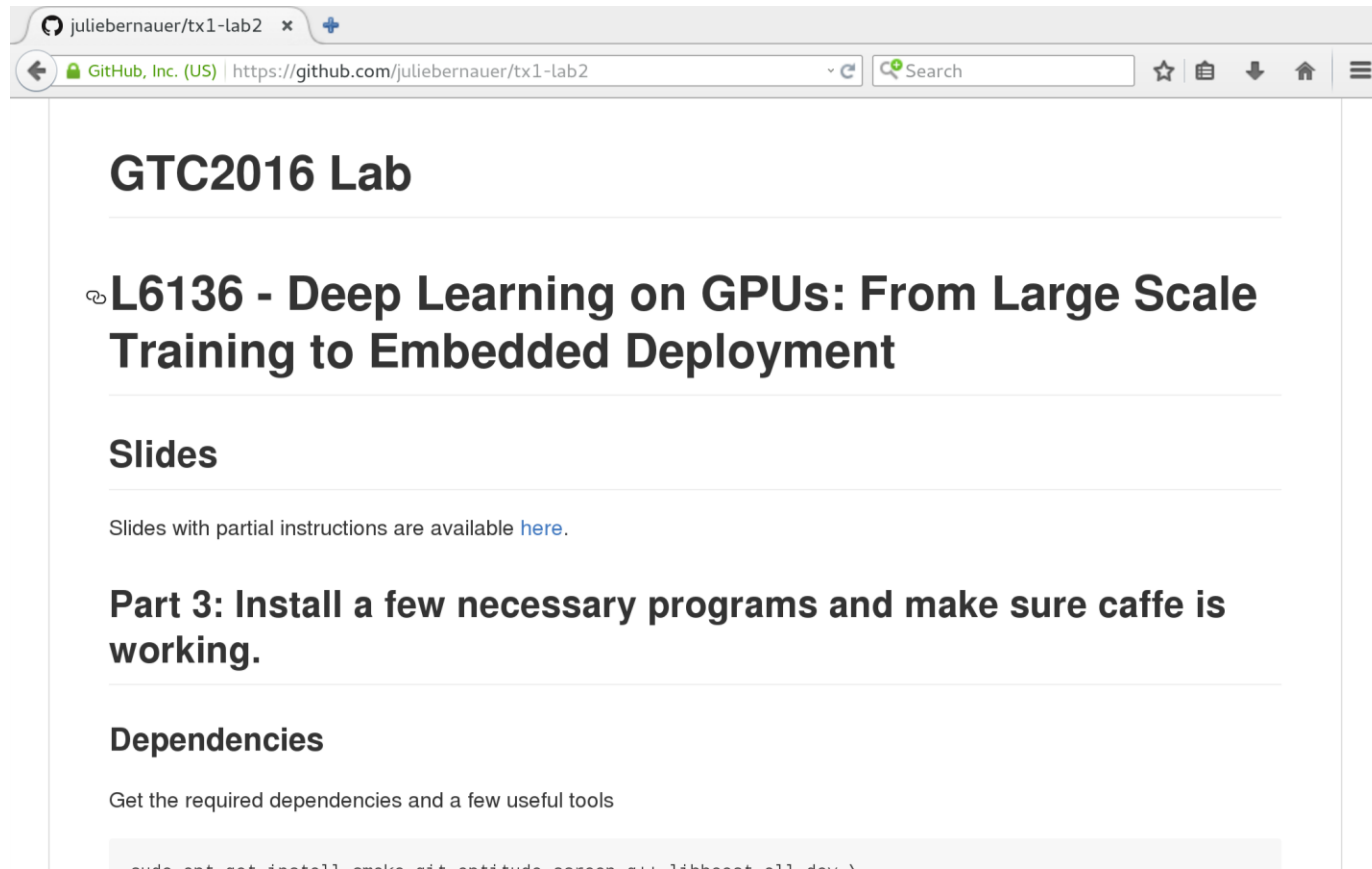
Scroll down to view your network's response

DEPLOYMENT ON TX1

3- CAFFE

TX1

Follow the instructions on
the Github page



The screenshot shows a web browser window displaying the GitHub repository page for 'juliebernauer/tx1-lab2'. The browser's address bar shows the URL 'https://github.com/juliebernauer/tx1-lab2'. The page content includes a section titled 'GTC2016 Lab' with a link to 'L6136 - Deep Learning on GPUs: From Large Scale Training to Embedded Deployment'. Below this, there is a 'Slides' section with a link to 'here'. The 'Part 3: Install a few necessary programs and make sure caffe is working.' section is also visible, followed by a 'Dependencies' section. The bottom of the page shows a code snippet starting with 'sudo apt-get install...'.

juliebernauer/tx1-lab2 x

GitHub, Inc. (US) | https://github.com/juliebernauer/tx1-lab2 Search

GTC2016 Lab

🔗 **L6136 - Deep Learning on GPUs: From Large Scale Training to Embedded Deployment**

Slides

Slides with partial instructions are available [here](#).

Part 3: Install a few necessary programs and make sure caffe is working.

Dependencies

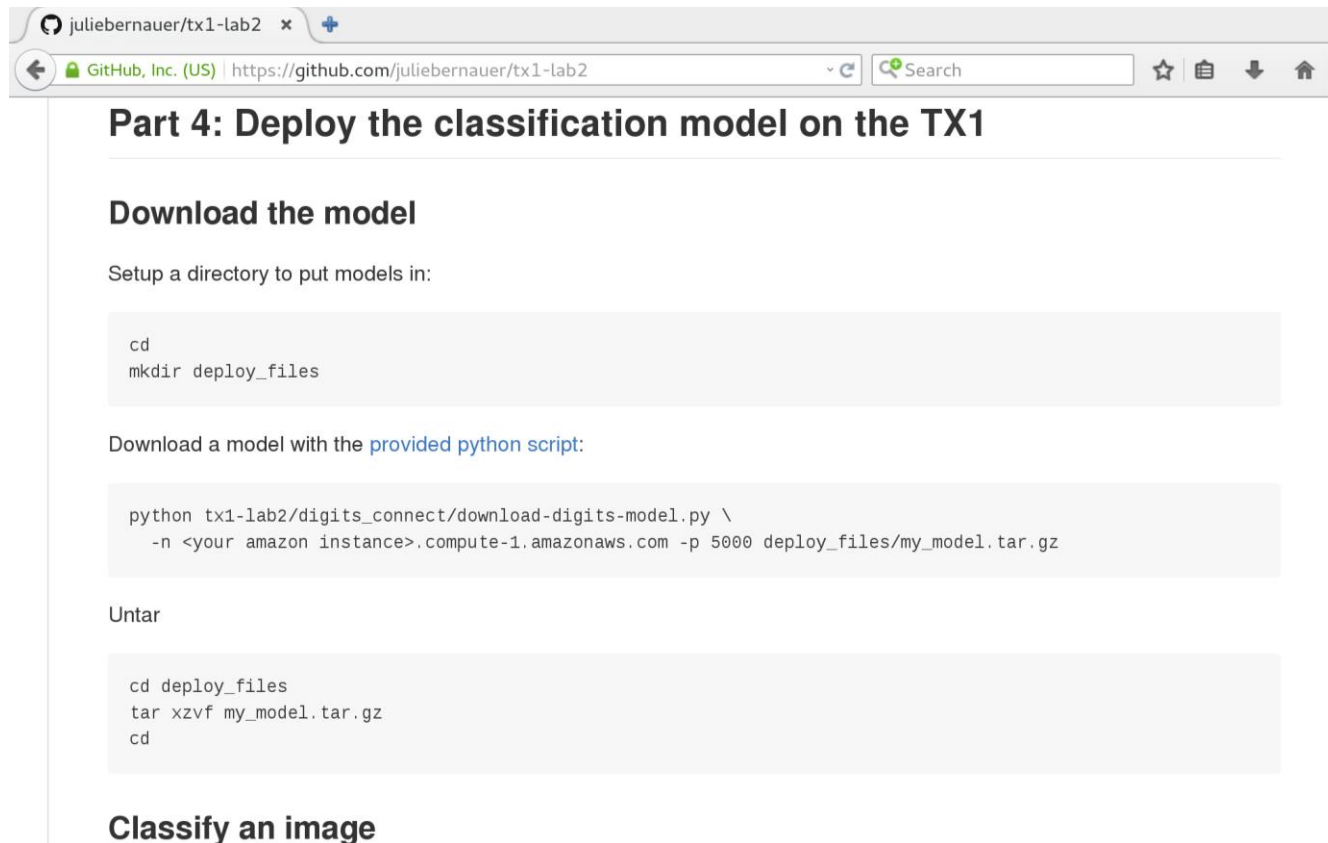
Get the required dependencies and a few useful tools

```
sudo apt-get install cmake git antitudo screen g++ libboost-all-dev \
```


4- DEPLOY THE IMAGE CLASSIFICATION MODEL

TX1

Follow the instructions on the Github page



The screenshot shows a web browser window with the address bar displaying "https://github.com/juliebernauer/tx1-lab2". The page title is "Part 4: Deploy the classification model on the TX1". The main content area is titled "Download the model" and includes the instruction "Setup a directory to put models in:" followed by a code block containing the commands "cd" and "mkdir deploy_files". Below this, it says "Download a model with the provided python script:" followed by a code block containing the command "python tx1-lab2/digits_connect/download-digits-model.py \ -n <your amazon instance>.compute-1.amazonaws.com -p 5000 deploy_files/my_model.tar.gz". The next section is titled "Untar" and includes a code block with the commands "cd deploy_files", "tar xzvf my_model.tar.gz", and "cd". The final section is titled "Classify an image".

Part 4: Deploy the classification model on the TX1

Download the model

Setup a directory to put models in:

```
cd
mkdir deploy_files
```

Download a model with the [provided python script](#):

```
python tx1-lab2/digits_connect/download-digits-model.py \
-n <your amazon instance>.compute-1.amazonaws.com -p 5000 deploy_files/my_model.tar.gz
```

Untar


```
cd deploy_files
tar xzvf my_model.tar.gz
cd
```

Classify an image

5- FUN WITH THE WEBCAM

TX1

Follow the instructions on
the Github page



GPU TECHNOLOGY
CONFERENCE

April 4-7, 2016 | Silicon Valley

THANK YOU

JOIN THE CONVERSATION

#GTC16   

JOIN THE NVIDIA DEVELOPER PROGRAM AT developer.nvidia.com/join

PRESENTED BY

